


Regression Tree

Typical Regression Tree Output

- In case of regression trees, the difference is that on the leaf nodes we have the means of the response variable values.



Regression Tree

- The data gets divided into two parts in the interest of decreasing the variation of response variable
- The child nodes have lesser variation than their respective parent nodes for response variable

Regression Tree in R

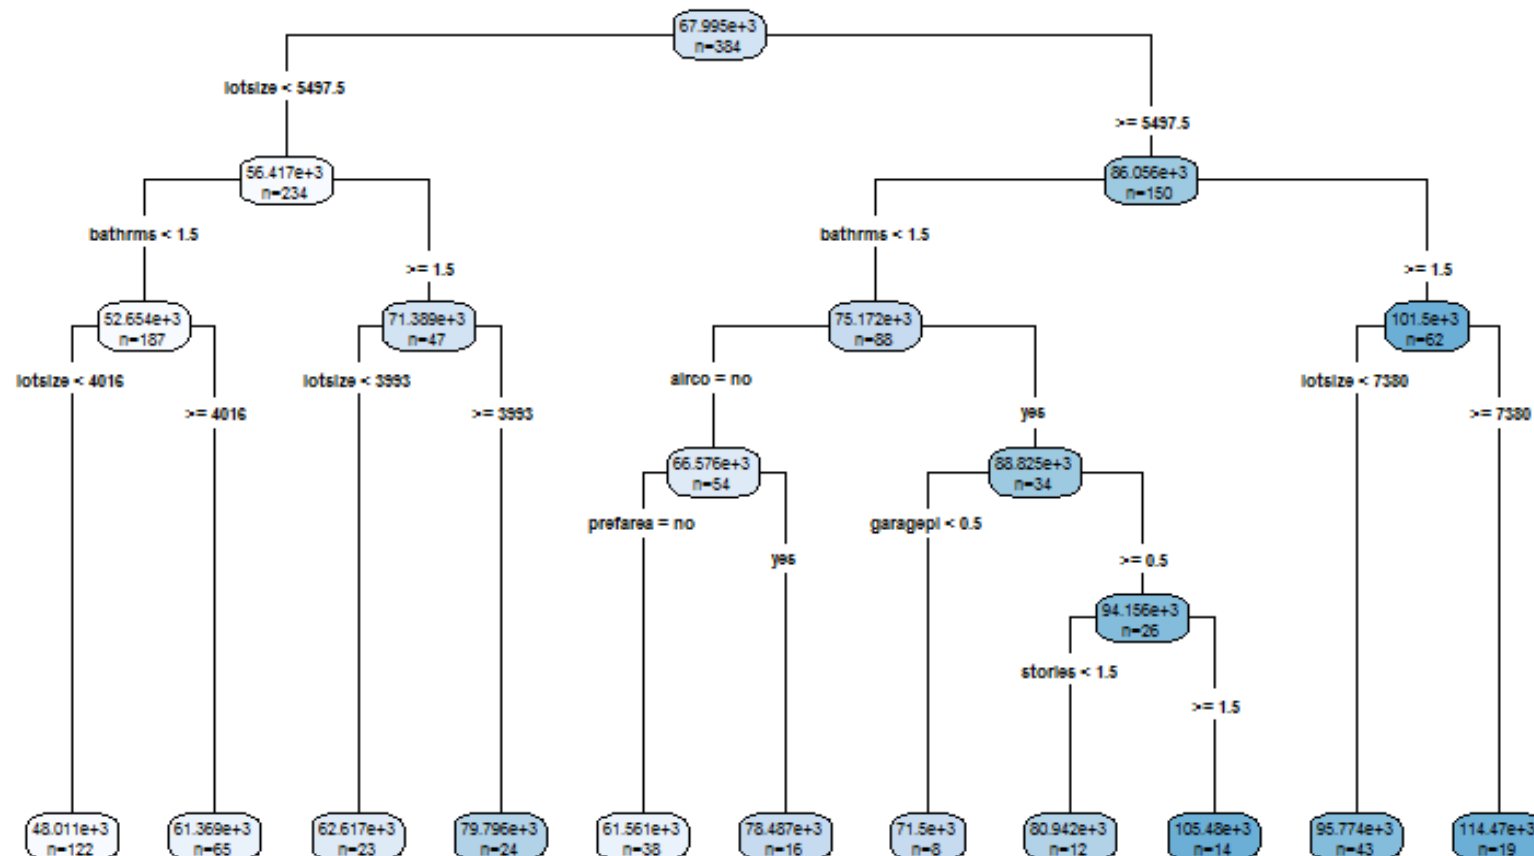
- For implementing Regression Tree in `rpart()`, option `method` should be specified with “`anova`”.
- For implementing Regression Tree in `ctree()`, there is no option to be specified differently. It identifies the response variable type and executes accordingly

Example : Sales Prices of Houses in the City of Windsor

- **Description**
 - a cross-section from 1987
 - *number of observations* : 546
 - *country* : Canada
- A dataframe containing :
 - **price** : sale price of a house
 - **lotsize** : the lot size of a property in square feet
 - **bedrooms** : number of bedrooms
 - **bathrms** : number of full bathrooms
 - **stories** : number of stories excluding basement
 - **driveway** : does the house has a driveway ?
 - **recroom** : does the house has a recreational room ?
 - **fullbase** : does the house has a full finished basement ?
 - **gashw** : does the house uses gas for hot water heating ?
 - **airco** : does the house has central air conditioning ?
 - **garagepl** : number of garage places
 - **prefarea** : is the house located in the preferred neighbourhood of the city ?

Program and Output – Using rpart

```
rpart.plot(fitRT,type = 4,extra = 1, digits = 5)
```



Model Evaluation: RMSE

- RMSE : Root Mean Square Error

$$RMSE = \sqrt{\frac{\sum (y_i - \hat{y}_i)^2}{n}}$$

where

y_i = Observed Values

\hat{y}_i = Predicted Values

n = No. of observations

Model Evaluation: MAPE

- MAPE: Mean Absolute Percentage Error

$$\text{MAPE} = \frac{\sum \left| \frac{y_i - \hat{y}_i}{y_i} \right|}{n}$$

- Where

y_i = Observed Values

\hat{y}_i = Predicted Values

n = No. of observations

Model Evaluation: RMSPE

- RMSPE: Root Mean Square Percentage Error
 - Often used at Kaggle competitions

$$\text{RMSPE} = \sqrt{\frac{1}{n} \sum_{i=1}^n \left(\frac{y_i - \hat{y}_i}{y_i} \right)^2}$$

- Where

y_i = Observed Values

\hat{y}_i = Predicted Values

n = No. of observations

Model Evaluation in R

- About MAPE, RMSE and RMPSE, only one criterion holds: **Smaller** their value, **Better** is the model prediction
- RMSE is calculated as directly by function `postResample()`
- For MAPE and RMSPE we can create a functions

```
MAPE <- function(y, yhat) {  
  mean(abs((y - yhat)/y))  
}
```


```
RMPSE<- function(y, yhat) {  
  sqrt(mean((y-yhat)/y)^2)  
}
```

Output

```
> pred.RT <- predict(fitRT,newdata = validation )
>
> postResample(pred.RT , validation$price)
      RMSE      Rsquared
1.809074e+04 5.414778e-01
>
> MAPE <- function(y, yhat) {
+   mean(abs((y - yhat)/y))
+ }
>
> MAPE(validation$price , pred.RT)
[1] 0.2036191
>
> RMSPE<- function(y, yhat) {
+   sqrt(mean((y-yhat)/y)^2)
+ }
>
> RMSPE(validation$price , pred.RT)
[1] 0.04037694
>
```

Pruning

- We take a look at `plotcp()` output

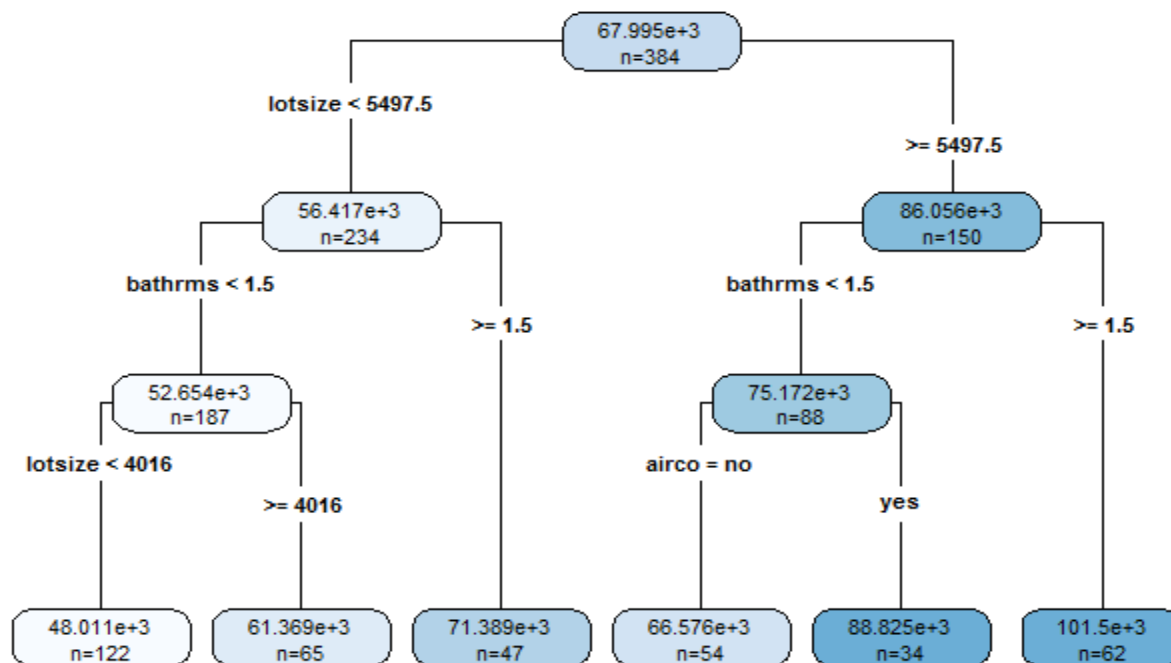


```
> fitRT$cpstable
```

	CP	nsplit	rel error	xerror	xstd
1	0.29348617	0	1.0000000	1.0051535	0.09724243
2	0.09217662	1	0.7065138	0.7250328	0.07373928
3	0.04818811	2	0.6143372	0.6572006	0.06866341
4	0.03775017	3	0.5661491	0.6367268	0.06719953
5	0.02765860	4	0.5283989	0.5957366	0.05874151
6	0.01683250	5	0.5007403	0.5945977	0.06134276
7	0.01285050	6	0.4839078	0.6042144	0.06289123
8	0.01266817	8	0.4582068	0.6093850	0.06347051
9	0.01179162	9	0.4455387	0.6086444	0.06344482
10	0.01000000	10	0.4337470	0.6039693	0.06348830

Pruning

```
fitRT.pruned <- prune(fitRT , cp=0.02765860 )
```



Output

```
> pred.RT.pruned <- predict(fitRT.pruned , newdata = validation)
> postResample(pred.RT.pruned , validation$price)
      RMSE      Rsquared
1.895032e+04 4.967799e-01
> MAPE(validation$price , pred.RT.pruned)
[1] 0.2180033
> RMSPE(validation$price , pred.RT.pruned)
[1] 0.04583555
```

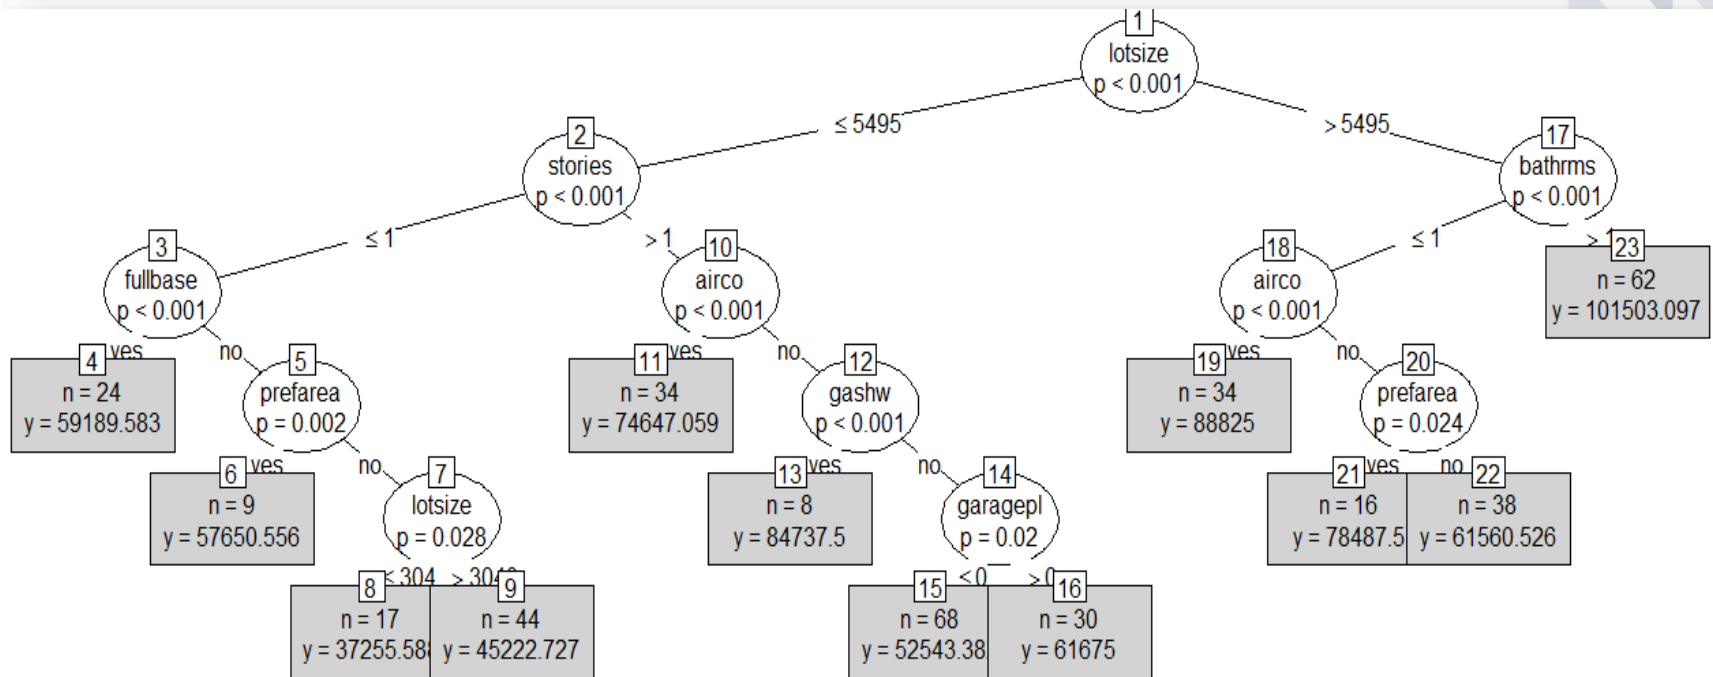
Though we might observe that the accuracy has been affected negatively by pruning, still we have minimized the risk of overfitting of the model

Program and Output – Using party

```
library(party)

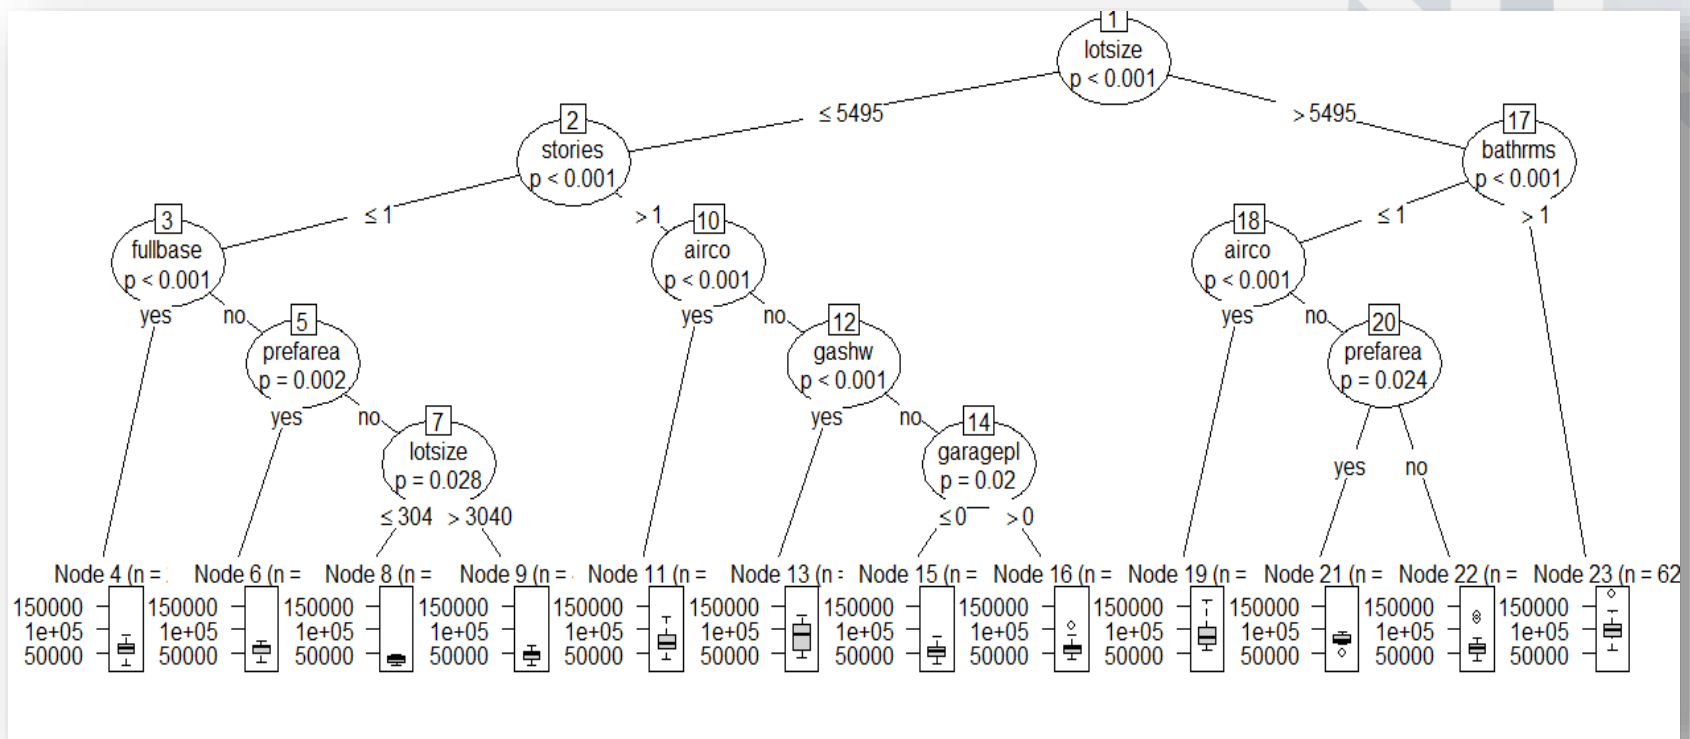
fitCT <- ctree(price ~ . , data = training )

plot(fitCT , type="simple")
```



Program and Output – Using party

```
plot(fitCT , type="extended" )
```



Accuracy Measures : party

```
> pred.CT <- predict(fitCT , newdata=validation)
>
> postResample(pred.CT , validation$price)
      RMSE      Rsquared
1.831758e+04 5.279237e-01
> MAPE(validation$price , pred.CT)
[1] 0.2145092
> RMSPE(validation$price , pred.CT)
[1] 0.06729213
```