

Random Forest

Ensembled Learning

- Random Forest is an ensembled learning approach
- In ensembled learning approach, multiple predictive models are developed and results are aggregated to improve the precision

Random Forest

- In this algorithm, the observations as well as variables are sampled to create multiple decision trees
- Each observation is classified by each decision tree.
- The outcome is considered as per the majority in different trees

Algorithm

(Considering N observations & M variables)

1. Sample out of N cases with replacement from the training set, many samples. Consider each sample as root node for the decision trees to be constructed.
2. Choose some $m < M$ number of variable by sampling at each node created in step 1.
3. Grow each tree without pruning with minimum node size as 1
4. Classify the validation / test set observations by traversing them through all the grown trees.
5. Classify each outcome by a majority vote of the trees.

OOB

- Each tree is constructed using a different bootstrap sample from the original data.
- About one-third of the cases are left out of the bootstrap sample and not used in the construction of the k th tree.
- Put each case left out in the construction of the k th tree, traverse the k th tree to get a classification.
- In this way, a test set classification is obtained for each case in about one-third of the trees. At the end of the run, take j to be the class that got most of the votes every time case n was oob.
- The proportion of times that j is not equal to the true class of n averaged over all cases is the oob error estimate. This has proven to be unbiased in many tests.
- This is done internally with the training set

Random Forest in R

- For Traditional Decision Trees (rpart type), we have function `randomForest()` package `randomForest`
- For conditional inference trees (party type), we have `cforest()` in package `party`

Package randomForest

- `randomForest()` in package `randomForest` implements Breiman's random forest algorithm

Syntax : `randomForest(formula, data, ntree, na.action, importance,...)`

Where

formula : Model formula

data : data frame object with train set

ntree : Number of trees to be grown

na.action : Action to be taken on NA values

importance : whether to create importance measure while creating randomForest object

Variable Importance

- The function `importance()` gives the list of variables along with the measure of their importance in terms of the purity gained by them.

Syntax : `importance(objRF , type, ...)`

Where

`objRF` : Object of class `randomForest`

`type` : either 1 or 2, specifying the type of importance measure (1=mean decrease in accuracy, 2=mean decrease in node impurity).

Example: Classification

- We will be using dataset churn from the package C50.
- The dataset has already been partitioned by its creators into two frames:
 - churnTrain : Training Set
 - churnTest : Validation / Test Set
- Intension is to predict the churning of the telecom customers

Attributes

- The dataset churn has the following attributes:
 - state(categorical)
 - account_length : how long account has been active
 - area_code
 - international_plan (yes/no)
 - voice_mail_plan (yes/no)
 - number_vmail_messages
 - total_day_minutes : minutes customer used service during the day
 - total_day_calls : total daily calls
 - total_day_charge : perhaps based on foregoing two variables

Attributes

- Attributes (Contd.):
 - total_eve_minutes : minutes customer used service during the evening
 - total_eve_calls : total evening calls
 - total_eve_charge : perhaps based on foregoing two variables
 - total_night_minutes : minutes customer used service during the night
 - total_night_calls : total night calls
 - total_night_charge : perhaps based on foregoing two variables
 - total_intl_minutes : minutes customer used service to make international calls
 - total_intl_calls : total international calls
 - total_intl_charge : perhaps based on foregoing two variables
 - number_customer_service_calls : Number of calls to customer service
 - churn (Response variable)

Program & Output

```
library(randomForest)
model.RF <- randomForest(churn ~ . , data = churnTrain ,
                          na.action=na.roughfix, importance=TRUE)
```

na.roughfix imputes
median/mode for missing
values

```
> model.RF
```

Call:

```
randomForest(formula = churn ~ . , data = churnTrain, importance = TRUE,      na.action
= na.roughfix)
```

Type of random forest: classification

Number of trees: 500

No. of variables tried at each split: 4

OOB estimate of error rate: 6.48%

Confusion matrix:

	yes	no	class.error
yes	397	86	0.17805383
no	130	2720	0.04561404

Importance

```
> importance(model.RF,type = 2)
```

	MeanDecreaseGini
state	110.643986
account_length	20.877809
area_code	4.566489
international_plan	56.297338
voice_mail_plan	14.256491
number_vmail_messages	21.713792
total_day_minutes	104.049412
total_day_calls	21.995356
total_day_charge	102.438042
total_eve_minutes	44.623016
total_eve_calls	19.105654
total_eve_charge	44.944149
total_night_minutes	25.375495
total_night_calls	20.275212
total_night_charge	24.307408
total_intl_minutes	30.511743
total_intl_calls	35.962697
total_intl_charge	30.884003
number_customer_service_calls	93.451117

Evaluation

```
> pred <- predict(model, newdata = churnTest[, -20])
```

```
>
```

```
> confusionMatrix( table(pred, churnTest$churn) )
```

Confusion Matrix and Statistics

pred	yes	no
yes	185	40
no	39	1403

Accuracy : 0.9526

95% CI : (0.9413, 0.9623)

No Information Rate : 0.8656

P-Value [Acc > NIR] : <2e-16

Kappa : 0.7967

Mcnemar's Test P-Value : 1

Sensitivity : 0.8259

Specificity : 0.9723

Pos Pred Value : 0.8222

Neg Pred Value : 0.9730

Prevalence : 0.1344

Detection Rate : 0.1110

Detection Prevalence : 0.1350

Balanced Accuracy : 0.8991

'Positive' Class : yes

Example: Numerical Response

- Consider the dataset Boston in package MASS with the following attributes:
 - crim: per capita crime rate by town.
 - zn: proportion of residential land zoned for lots over 25,000 sq.ft.
 - indus: proportion of non-retail business acres per town.
 - chas: Charles River dummy variable (= 1 if tract bounds river; 0 otherwise).
 - nox: nitrogen oxides concentration (parts per 10 million).
 - rm: average number of rooms per dwelling.
 - age: proportion of owner-occupied units built prior to 1940.
 - dis: weighted mean of distances to five Boston employment centres.
 - rad: index of accessibility to radial highways.
 - tax: full-value property-tax rate per \backslash \$10,000.
 - ptratio: pupil-teacher ratio by town.
 - black: $1000(B_k - 0.63)^2$ where B_k is the proportion of blacks by town.
 - lstat: lower status of the population (percent).
 - medv: median value of owner-occupied homes in \backslash \$1000s.

Program & Output

```
library(randomForest)
model.RF <- randomForest(medv ~ . , data = training ,
                          na.action=na.roughfix, importance=TRUE)
```

```
> model.RF
```

Call:

```
randomForest(formula = medv ~ ., data = training, importance = TRUE,      na.action =
na.roughfix)
```

 Type of random forest: regression

 Number of trees: 500

No. of variables tried at each split: 4

 Mean of squared residuals: 12.3705

 % Var explained: 86.08

Variable Importance

```
> importance(model.RF,type = 2)
```

	IncNodePurity
crim	1873.85795
zn	153.04752
indus	1661.92180
chas	85.53658
nox	2249.01481
rm	9214.13244
age	808.49074
dis	1773.60915
rad	295.72698
tax	1221.70373
ptratio	1876.61594
black	647.32592
lstat	9258.32671

Evaluation

```
> pred.RF <- predict(model.RF, newdata = validation[, -14])
>
> postResample(pred.RF, validation$medv)
      RMSE  Rsquared
2.9496507 0.8871838
>
> MAPE <- function(y, yhat) {
+   mean(abs((y - yhat)/y))
+ }
>
> MAPE(validation$medv, pred.RF)
[1] 0.1020432
>
> RMSPE <- function(y, yhat) {
+   sqrt(mean((y - yhat)/y)^2)
+ }
>
> RMSPE(validation$medv, pred.RF)
[1] 0.04299174
```

Using cforest()

```
library(party)
model.RF <- cforest(medv ~ ., data = Boston,
                    control = cforest_unbiased(ntree = 50))
```

```
> pred.RF <- predict(model.RF, newdata = validation[,-14])
>
> postResample(pred.RF , validation$medv)
      RMSE  Rsquared
2.9072959 0.8913023
>
> MAPE <- function(y, yhat) {
+   mean(abs((y - yhat)/y))
+ }
>
> MAPE(validation$medv , pred.RF)
[1] 0.09490211
>
> RMSPE<- function(y, yhat) {
+   sqrt(mean((y-yhat)/y)^2)
+ }
>
> RMSPE(validation$medv , pred.RF)
[1] 0.04144711
```