

Smoothing Methods

Example Dataset

- We will use a time series dataset of Monthly Milk Production (in pounds) of cows from January 1962 to December 1975
- Source: <http://data.is/1qY3LDd>
- Website: www.datamarket.com

```
mlk <- read.csv("monthly-milk-production-pounds-p.csv")  
milk_ts <- ts(data = mlk$Milk, start = c(1962,1), frequency = 12 )
```

Data Partition

- We have divided the data into
 - Training Data : Data from Jan 1962 to November 1972
 - Validation Data : Data from December 1972 to December 1975

```
# Number of Observations in Validation data
nValid <- 38
# Number of Observations in Training data
nTrain <- nrow(milk) - nValid
# Training data and Validation data partitioned using window()
train.ts <- window(milk_ts, start = c(1962,1), end = c(1962, nTrain + 1))
valid.ts <- window(milk_ts, start = c(1962, nTrain + 2))
```

Smoothing Methods

- Smoothing Methods are a kind of forecasting methods that are data driven
- These methods directly estimate time series components from the data
- We will be learning:
 - Moving Average
 - Simple Smoothing
 - Holt's Method
 - Holt-Winter's Method

Moving Average

- The consecutive values of the time series are averaged with a specific width maintained.
- A moving average with width **w** means average taken across each set of **w** consecutive time series values, where **w** is an integer input by the user.
- There are two types of moving averages:
 - Centered Moving Average
 - Trailing Moving Average

Centered Moving Average

- Centered Moving Average are powerful for visualization
- The value of the moving average at time t is computed by centering the time span around time t and averaging across w values within the time span
- The goal is to suppress the seasonality to better visualize the trend. Hence choosing width as length of seasonal cycle is more desirable

Centered MA Calculations

- With a time span $w=5$, the moving average at time point $t=3$ would be average of 1st, 2nd, 3rd, 4th, 5th time points.
- At time span $w=4$, moving average would be average of 2nd, 3rd, 4th, 5th, 6th time points
- Function **ma()** in the *forecast* package creates a moving average

Syntax : `ma(x, order, center=TRUE)`

Where

`x` : Object of class `ts`

`order` : span of moving average

`center` : logical value, if true then moving average is centered for even orders


When w is even

- When order w is even then, centered MA is calculated as average of the two asymmetric moving averages
- When order $w = 4$,

$$MA_t = \frac{\left[\frac{(y_{t-2} + y_{t-1} + y_t + y_{t+1})}{4} + \frac{(y_{t-1} + y_t + y_{t+1} + y_{t+2})}{4} \right]}{2}$$

Centered MA Example

```
library(forecast)
maMilk1Cent <- ma(milk_ts, order = 3)
plot(milk_ts, ylab = "Milk Production", xlab = "Time", bty = "l", xaxt = "n", main = "")
axis(1, at = seq(1962, 1975, 1), labels = format(seq(1962, 1975, 1)))
lines(maMilk1Cent, lwd = 2, col = "blue")
```



Trailing Moving Average

- Centered MAs use both past and future time points, so they cannot be used for forecasting
- For forecasting trailing moving average can be used because here average is calculated in a time span for past and most recent time point
- The k-step ahead forecast F_{t+k} is computed with the formula:

$$F_{t+k} = (y_t + y_{t-1} + \dots + y_{t-w+1}) / w$$

Trailing MA Calculation

- Trailing MA can be calculated with the function `rollmean()` in package `zoo`.

Syntax : `rollmean(ts , w , align, ...)`

Where


`ts` : Time series Object

`w` : Window span for moving average

`align` : character specifying whether the index of the result should be left- or right-aligned or centered

Trailing MA Example

```
library(zoo)
maMilk3Trail <- rollmean(train.ts, k = 3, align = "right")
lastMA <- tail(maMilk3Trail, 1)
predMA <- ts(rep(lastMA, nvalid), start = c(1962, nTrain + 1),
             end = c(1962, nTrain + nvalid), freq = 12)
plot(train.ts, ylab = "Milk Production", xlab = "Time", bty = "l",
     xaxt = "n", main = "")
axis(1, at = seq(1962, 1975, 1), labels = format(seq(1962, 1975, 1)))
lines(maMilk3Trail, lwd = 2, col = "blue")
lines(predMA, lwd = 2, col = "blue", lty = 2)
lines(valid.ts)
```



Accuracy

- The function `accuracy()` calculates various summary measures of forecast accuracy

Syntax : `accuracy(f , a , ...)`

where

`f` : an object of class `forecast` or a numerical vector containing the forecasted values

`a` : an object of class `forecast` or a numerical vector containing the actual values

Measures of Accuracy

The measures calculated are:

- ME: Mean Error
- RMSE: Root Mean Squared Error
- MAE: Mean Absolute Error
- MPE: Mean Percentage Error
- **MAPE: Mean Absolute Percentage Error**
- **MASE: Mean Absolute Scaled Error**
- ACF1: Autocorrelation of errors at lag 1.
- Theil's U: Uncertainty Coefficient

```
> accuracy(predMA,valid.ts)
```

	ME	RMSE	MAE	MPE	MAPE	ACF1	Theil's U
Test set	58.21622	83.45836	64.91892	6.371956	7.239719	0.6635945	1.582979

Differencing


- Differencing means taking the difference between values of time points
- Differencing may be between consecutive values or between values with lag more than one
- A lag-1 difference means taking the difference between two consecutive time points ($y_t - y_{t-1}$)
- A lag-k difference means taking difference between value and its value k-periods back ($y_t - y_{t-k}$)

Removing Trend and De-seasonalizing

- Lag-1 differencing is useful for removing trend
- To remove seasonal pattern with M seasons, we can difference at lag M
- e.g. To remove a day-of-week pattern in daily data, we can take lag-7 differences

Differencing Example

```
df.train.1 <- diff(train.ts, lag = 1)
library(zoo)
maMilk3Trail <- rollmean(df.train.1, k = 3, align = "right")
lastMA <- tail(maMilk3Trail, 1)
predMA <- ts(rep(lastMA, nValid), start = c(1962, nTrain + 1),
             end = c(1962, nTrain + nValid), freq = 12)
plot(df.train.1, ylab = "Milk Production", xlab = "Time", bty = "l",
     xaxt = "n", main = "")
axis(1, at = seq(1962, 1975, 1), labels = format(seq(1962, 1975, 1)))
lines(maMilk3Trail, lwd = 2, col = "blue")
lines(predMA, lwd = 2, col = "blue", lty = 2)
```



Simple Exponential Smoothing

- In Simple Exponential Smoothing, weighted average of all past values is taken in such a way that the weights decrease exponentially into past
- Like Moving Average, this method is used for forecasting series that have no trend and no seasonality

Calculation

- The exponential smoother calculates a forecast at time $t+1$, F_{t+1} :

$$F_{t+1} = \alpha y_t + \alpha (1 - \alpha)y_{t-1} + \alpha (1 - \alpha)^2 y_{t-2} + \dots$$

- Where α is a constant between 0 and 1 called smoothing constant
- The above equation can also be written as:
$$F_{t+1} = F_t + \alpha e_t$$
 - Where F_t is forecast at time t and e_t is forecast error at time t

Choice of α

- The smoothing constant α determines the rate of learning
- A value close to 1 implies fast learning, i.e. the most recent values influence the forecast most
- A value close to 0 implies slow learning, i.e. the past observations influence the forecast most
- The default values of α that have been mostly observed to work well are between 0.1 and 0.2.

Simple Smoothing in R

- We can calculate time series estimates of simple exponential smoothing with the help of **ses()** in package ***forecast***.

Syntax : `ses(ts , h , alpha)`

where


ts : a numeric vector or time series

h : Number of periods for forecasting

alpha : Smoothing constant for level, if not specified
it will be estimated by the function

Simple Exponential Example

```
sesMilk <- ses(train.ts, h = nvalid, alpha = 0.2)
plot(sesMilk, ylab = "Milk Production", xlab = "Time", bty = "l",
     xaxt = "n", main = "", flty = 2)
axis(1, at = seq(1962, 1975, 1), labels = format(seq(1962, 1975, 1)))
lines(sesMilk$fitted, lwd = 2, col = "blue")
lines(valid.ts)
```



Model & Accuracy

```
> sesMilk$model
Simple exponential smoothing

Call:
ses(x = train.ts, h = nValid, alpha = 0.2)

Smoothing parameters:
  alpha = 0.2

Initial states:
  l = 620.6362

sigma: 61.2608

      AIC      AICc      BIC
1718.818 1718.849 1721.693
```

```
accuracy(sesMilk , valid.ts)
```

	ME	RMSE	MAE	MPE	MAPE	MASE	ACF1	Theil's U
Training set	8.198829	61.26084	51.90699	0.5052553	7.075074	2.153742	0.6857870	NA
Test set	17.770688	62.38563	51.60958	1.6086402	5.907381	2.141402	0.6635945	1.194783

ses() without alpha

```
> sesMilk_opt <- ses(train.ts, h = nValid)
> plot(sesMilk_opt, ylab = "Milk Production", xlab = "Time", bty = "l", xaxt = "n", main = "", flty = 2)
> axis(1, at = seq(1962, 1975, 1), labels = format(seq(1962, 1975, 1)))
> lines(sesMilk_opt$fitted, lwd = 2, col = "blue")
> lines(valid.ts)
> sesMilk_opt$model
```

Simple exponential smoothing

Call:

```
ses(x = train.ts, h = nValid)
```

Smoothing parameters:

```
alpha = 0.9999
```

Initial states:

```
l = 589.0255
```

```
sigma: 44.1128
```

	AIC	AICc	BIC
	1634.779	1634.873	1640.530

```
> accuracy(sesMilk_opt, valid.ts)
```

	ME	RMSE	MAE	MPE	MAPE	MASE	ACF1	Theil's U
Training set	1.351119	44.11282	38.00845	0.01697799	5.209431	1.577059	0.06249857	NA
Test set	87.211813	105.74531	87.53638	9.78679961	9.829505	3.632088	0.66359455	2.017068

ses() without alpha

